
Formatting Instructions For NeurIPS 2026

Anonymous Author(s)

Affiliation

Address

email

1 Neur

Abstract

2 The abstract paragraph should be indented 1/2 inch (3 picas) on both the left- and
3 right-hand margins. Use 10 point type, with a vertical spacing (leading) of 11 points.
4 The word **Abstract** must be centered, bold, and in point size 12. Two line spaces
5 precede the abstract. The abstract must be limited to one paragraph.

6 1 Submission of papers to NeurIPS 2026

7 Please read the instructions below carefully and follow them faithfully.

8 1.1 Style

9 Papers to be submitted to NeurIPS 2026 must be prepared according to the instructions presented here.
10 Papers may only be up to **nine** pages long, including figures. **Papers that exceed the page limit will**
11 **not be reviewed (or in any other way considered) for presentation at the conference.** Additional
12 pages *containing acknowledgments, references, checklist, and optional technical appendices* do not
13 count as content pages.

14 The margins in 2026 are the same as those in previous years.

15 Authors are required to use the NeurIPS L^AT_EX style files obtainable at the NeurIPS website as
16 indicated below. Please make sure you use the current files and not previous versions. Tweaking the
17 style files may be grounds for desk rejection.

18 1.2 Retrieval of style files

19 The style files for NeurIPS and other conference information are available on the website at

20 <https://neurips.cc>.

21 The only supported style file for NeurIPS 2026 is `neurips_2026.sty`, rewritten for L^AT_EX 2_ε.
22 **Previous style files for L^AT_EX 2.09, Microsoft Word, and RTF are no longer supported.**

23 The L^AT_EX style file contains three optional arguments:

- 24 • `final`, which creates a camera-ready copy,
- 25 • `preprint`, which creates a preprint for submission to, e.g., arXiv,
- 26 • `nonatbib`, which will not load the `natbib` package for you in case of package clash.

27 **Preprint option** If you wish to post a preprint of your work online, e.g., on arXiv, using the
28 NeurIPS style, please use the `preprint` option. This will create a nonanonymized version of your

work with the text “Preprint. Work in progress.” in the footer. This version may be distributed as you see fit, as long as you do not say which conference it was submitted to. Please **do not** use the `final` option, which should **only** be used for papers accepted to NeurIPS.

At submission time, please omit the `final` and `preprint` options. This will anonymize your submission and add line numbers to aid review. Please do *not* refer to these line numbers in your paper as they will be removed during generation of camera-ready copies.

The file `neurips_2026.tex` may be used as a “shell” for writing your paper. All you have to do is replace the author, title, abstract, and text of the paper with your own.

The formatting instructions contained in these style files are summarized in Sections ??, ??, and ?? below.

2 General formatting instructions

The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long. The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing (leading) of 11 points. Times New Roman is the preferred typeface throughout, and will be selected for you by default. Paragraphs are separated by $\frac{1}{2}$ line space (5.5 points), with no indentation.

The paper title should be 17 point, initial caps/lower case, bold, centered between two horizontal rules. The top rule should be 4 points thick and the bottom rule should be 1 point thick. Allow $\frac{1}{4}$ inch space above and below the title to rules. All pages should start at 1 inch (6 picas) from the top of the page.

For the final version, authors’ names are set in boldface, and each name is centered above the corresponding address. The lead author’s name is to be listed first (left-most), and the co-authors’ names (if different address) are set to follow. If there is only one co-author, list both author and co-author side by side.

Please pay special attention to the instructions in Section ?? regarding figures, tables, acknowledgments, and references.

3 Headings: first level

All headings should be lower case (except for first word and proper nouns), flush left, and bold.

First-level headings should be in 12-point type.

3.1 Headings: second level

Second-level headings should be in 10-point type.

3.1.1 Headings: third level

Third-level headings should be in 10-point type.

Paragraphs There is also a `\paragraph` command available, which sets the heading in bold, flush left, and inline with the text, with the heading followed by 1 em of space.

4 Citations, figures, tables, references

These instructions apply to everyone.

4.1 Citations within the text

The `natbib` package will be loaded for you by default. Citations may be author/year or numeric, as long as you maintain internal consistency. As to the format of the references themselves, any style is acceptable as long as it is used consistently.




Figure 1: Sample figure caption. Explain what the figure shows and add a key take-away message to the caption.

69 The documentation for natbib may be found at

70 `http://mirrors.ctan.org/macros/latex/contrib/natbib/natnotes.pdf`

71 Of note is the command `\citet`, which produces citations appropriate for use in inline text. For
72 example,

73 `\citet{hasselmo}` investigated\dotso

74 produces

75 Hasselmo, et al. (1995) investigated...

76 If you wish to load the natbib package with options, you may add the following before loading the
77 neurips_2026 package:

78 `\PassOptionsToPackage{options}{natbib}`

79 If natbib clashes with another package you load, you can add the optional argument nonatbib
80 when loading the style file:

81 `\usepackage[nonatbib]{neurips_2026}`

82 As submission is double blind, refer to your own published work in the third person. That is, use “In
83 the previous work of Jones et al. [4],” not “In our previous work [4].” If you cite your other papers
84 that are not widely available (e.g., a journal paper under review), use anonymous author names in the
85 citation, e.g., an author of the form “A. Anonymous” and include a copy of the anonymized paper in
86 the supplementary material.

87 4.2 Footnotes

88 Footnotes should be used sparingly. If you do require a footnote, indicate footnotes with a number¹
89 in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote
90 with a horizontal rule of 2 inches (12 picas).

91 Note that footnotes are properly typeset *after* punctuation marks.²

92 4.3 Figures

93 All artwork must be neat, clean, and legible. Lines should be dark enough for reproduction purposes.
94 The figure number and caption always appear after the figure. Place one line space before the figure
95 caption and one line space after the figure. The figure caption should be lower case (except for the
96 first word and proper nouns); figures are numbered consecutively.

¹Sample of the first footnote.

²As in this example.

Table 1: Sample table caption. Explain what the table shows and add a key take-away message to the caption.

Part		
Name	Description	Size (μm)
Dendrite	Input terminal	≈ 100
Axon	Output terminal	≈ 10
Soma	Cell body	up to 10^6

97 You may use color figures. However, it is best for the figure captions and the paper body to be legible
98 if the paper is printed in either black/white or in color.

99 4.4 Tables

100 All tables must be centered, neat, clean, and legible. The table number and title always appear before
101 the table. See Table ??.

102 Place one line space before the table title, one line space after the table title, and one line space after
103 the table. The table title must be lower case (except for the first word and proper nouns); tables are
104 numbered consecutively.

105 Note that publication-quality tables *do not contain vertical rules*. We strongly suggest the use of the
106 booktabs package, which allows for typesetting high-quality, professional tables:

107 <https://www.ctan.org/pkg/booktabs>

108 This package was used to typeset Table ??.

109 4.5 Math

110 Note that display math in bare TeX commands will not create correct line numbers for sub-
111 mission. Please use LaTeX (or AMSTeX) commands for unnumbered display math. (You
112 really shouldn't be using \$\$ anyway; see <https://tex.stackexchange.com/questions/503/why-is-preferable-to> and [https://tex.stackexchange.com/questions/40492/](https://tex.stackexchange.com/questions/40492/what-are-the-differences-between-align-equation-and-displaymath)
113 [what-are-the-differences-between-align-equation-and-displaymath](https://tex.stackexchange.com/questions/40492/what-are-the-differences-between-align-equation-and-displaymath) for more infor-
114 mation.)
115

116 4.6 Final instructions

117 Do not change any aspects of the formatting parameters in the style files. In particular, do not modify
118 the width or length of the rectangle the text should fit into, and do not change font sizes. Please note
119 that pages should be numbered.

120 5 Preparing PDF files

121 Please prepare submission files with paper size "US Letter," and not, for example, "A4."

122 Fonts were the main cause of problems in the past years. Your PDF file must only contain Type 1 or
123 Embedded TrueType fonts. Here are a few instructions to achieve this.

- 124 • You should directly generate PDF files using pdf_latex.
- 125 • You can check which fonts a PDF files uses. In Acrobat Reader, select the menu
126 Files>Document Properties>Fonts and select Show All Fonts. You can also use the program
127 pdf_fonts which comes with xpdf and is available out-of-the-box on most Linux machines.
- 128 • xfig "patterned" shapes are implemented with bitmap fonts. Use "solid" shapes instead.
- 129 • The \b_bold package almost always uses bitmap fonts. You should use the equivalent AMS
130 Fonts:

131 `\usepackage{amsfonts}`
 132 followed by, e.g., `\mathbb{R}`, `\mathbb{N}`, or `\mathbb{C}` for $\mathbb{R}$, $\mathbb{N}$ or $\mathbb{C}$. You can also
 133 use the following workaround for reals, natural and complex:
 134 `\newcommand{\RR}{I\!\!R} %real numbers`
 135 `\newcommand{\Nat}{I\!\!N} %natural numbers`
 136 `\newcommand{\CC}{I\!\!C} %complex numbers`
 137 Note that `amsfonts` is automatically loaded by the `amssymb` package.

138 If your file contains type 3 fonts or non embedded TrueType fonts, we will ask you to fix it.

139 5.1 Margins in L^AT_EX

140 Most of the margin problems come from figures positioned by hand using `\special` or other
 141 commands. We suggest using the command `\includegraphics` from the `graphicx` package.
 142 Always specify the figure width as a multiple of the line width as in the example below:

143 `\usepackage[pdftex]{graphicx} ...`
 144 `\includegraphics[width=0.8\linewidth]{myfile.pdf}`

145 See Section 4.4 in the graphics bundle documentation ([http://mirrors.ctan.org/macros/](http://mirrors.ctan.org/macros/latex/required/graphics/grfguide.pdf)
 146 [latex/required/graphics/grfguide.pdf](http://mirrors.ctan.org/macros/latex/required/graphics/grfguide.pdf))

147 A number of width problems arise when L^AT_EX cannot properly hyphenate a line. Please give LaTeX
 148 hyphenation hints using the `\-` command when necessary.

149 References

150 References follow the acknowledgments in the camera-ready paper. Use unnumbered first-level
 151 heading for the references. Any choice of citation style is acceptable as long as you are consistent. It
 152 is permissible to reduce the font size to `small` (9 point) when listing the references. Note that the
 153 Reference section does not count towards the page limit.

154 [1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In
 155 G. Tesauro, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp.
 156 609–616. Cambridge, MA: MIT Press.

157 [2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the*
 158 *GENeral NEural Simulation System*. New York: TELOS/Springer-Verlag.

159 [3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent
 160 synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.

161 A Technical appendices and supplementary material

162 Technical appendices with additional results, figures, graphs, and proofs may be submitted with the
 163 paper submission before the full submission deadline (see above). You can upload a ZIP file for
 164 videos or code, but do not upload a separate PDF file for the appendix. There is no page limit for the
 165 technical appendices.

166 Note: Think of the appendix as “optional reading” for reviewers. The paper must be able to stand
 167 alone without the appendix; for example, adding critical experiments that support the main claims to
 168 an appendix is inappropriate.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract, Introduction, and Table ?? separate established framework evidence from the unfinished ranking outcome.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 discusses baseline provenance, lineage drift, seeds, validation overfitting, dependence, external validity, and the security boundary.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper introduces notation for the measurement protocol but contains no theorem or proof claim.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: Appendix ?? reproduces the focused framework test, but the external cluster artifacts and final certified campaign are not yet released.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

275 Question: Does the paper provide open access to the data and code, with sufficient instruc-
276 tions to faithfully reproduce the main experimental results, as described in supplemental
277 material?

278 Answer: [No]

279 Justification: The framework source is documented in the working repository, but a per-
280 manent public artifact and redistributable data/feature package are not yet cited in this
281 version.

282 Guidelines:

- 283 • The answer [N/A] means that paper does not include experiments requiring code.
- 284 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/](https://neurips.cc/public/guides/CodeSubmissionPolicy)
285 [public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 286 • While we encourage the release of code and data, we understand that this might not
287 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
288 including code, unless this is central to the contribution (e.g., for a new open-source
289 benchmark).
- 290 • The instructions should contain the exact command and environment needed to run to
291 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 292 • The authors should provide instructions on data access and preparation, including how
293 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 294 • The authors should provide scripts to reproduce all experimental results for the new
295 proposed method and baselines. If only a subset of experiments are reproducible, they
296 should state which ones are omitted from the script and why.
- 297 • At submission time, to preserve anonymity, the authors should release anonymized
298 versions (if applicable).
- 299 • Providing as much information as possible in supplemental material (appended to the
300 paper) is recommended, but including URLs to data and code is permitted.

302 6. Experimental setting/details

303 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
304 rameters, how they were chosen, type of optimizer) necessary to understand the results?

305 Answer: [No]

306 Justification: Section 4 and Appendix ?? specify the intended design, but corrected recipes,
307 seed schedule, frozen lineage rules, and final resource settings remain incomplete.

308 Guidelines:

- 309 • The answer [N/A] means that the paper does not include experiments.
- 310 • The experimental setting should be presented in the core of the paper to a level of detail
311 that is necessary to appreciate the results and make sense of them.
- 312 • The full details can be provided either with the code, in appendix, or as supplemental
313 material.

314 7. Experiment statistical significance

315 Question: Does the paper report error bars suitably and correctly defined or other appropriate
316 information about the statistical significance of the experiments?

317 Answer: [No]

318 Justification: No inferential claim is made. Multi-seed and dependence-aware uncertainty
319 are explicitly required before the final ranking claim.

320 Guidelines:

- 321 • The answer [N/A] means that the paper does not include experiments.
- 322 • The authors should answer [Yes] if the results are accompanied by error bars, confidence
323 intervals, or statistical significance tests, at least for the experiments that support the
324 main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The focused CPU test time is reported, but the cluster campaign’s complete accelerator, wall-clock, failed-run, and LLM-cost ledger is not available locally.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a research-operations and retrospective benchmark study; the paper avoids clinical-use claims and states its evidence boundaries.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Broader Impact section covers transparency benefits, premature-leaderboard risk, patient-data handling, and non-clinical scope.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: No high-risk model, scraped dataset, or generative media asset is released by this paper.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The papers and cohort sources are credited, but asset-by-asset versions and license terms must be added with the final public artifact.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: and are new code assets, but the permanent release, version, license manifest, and complete campaign documentation are not yet provided in the paper.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study uses no crowdsourcing and recruits no participants.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper reports no new intervention or participant recruitment; final secondary-use governance language must accompany the frozen cohort manifests.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: The core method uses coding agents, and Appendix ?? plus the Author and Tooling Disclosure describe their methodological and writing roles.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.